

# **CENG 311**

## **Computer Architecture**

### **Lecture 9**

1

## **8086 Assembly on Linux**

*Asst. Prof. Tolga Ayav, Ph.D.*

Department of Computer Engineering  
İzmir Institute of Technology

# Assemblers

- as, as86, gas, nasm ...
- We use nasm:
  - Portable: DOS, Linux, Windows versions
  - Free, widely used
- Differences between gas and nasm:
  - as, gas use AT&T syntax: % prefixes to register names.
  - Nasm use Intel syntax
  - In Nasm: Destination first, source operand second
  - In as, gas: Source operand first, destination second
- We will be using nasm but be careful when you use inline assembly in gcc, it uses AT&T syntax.

# Assembling

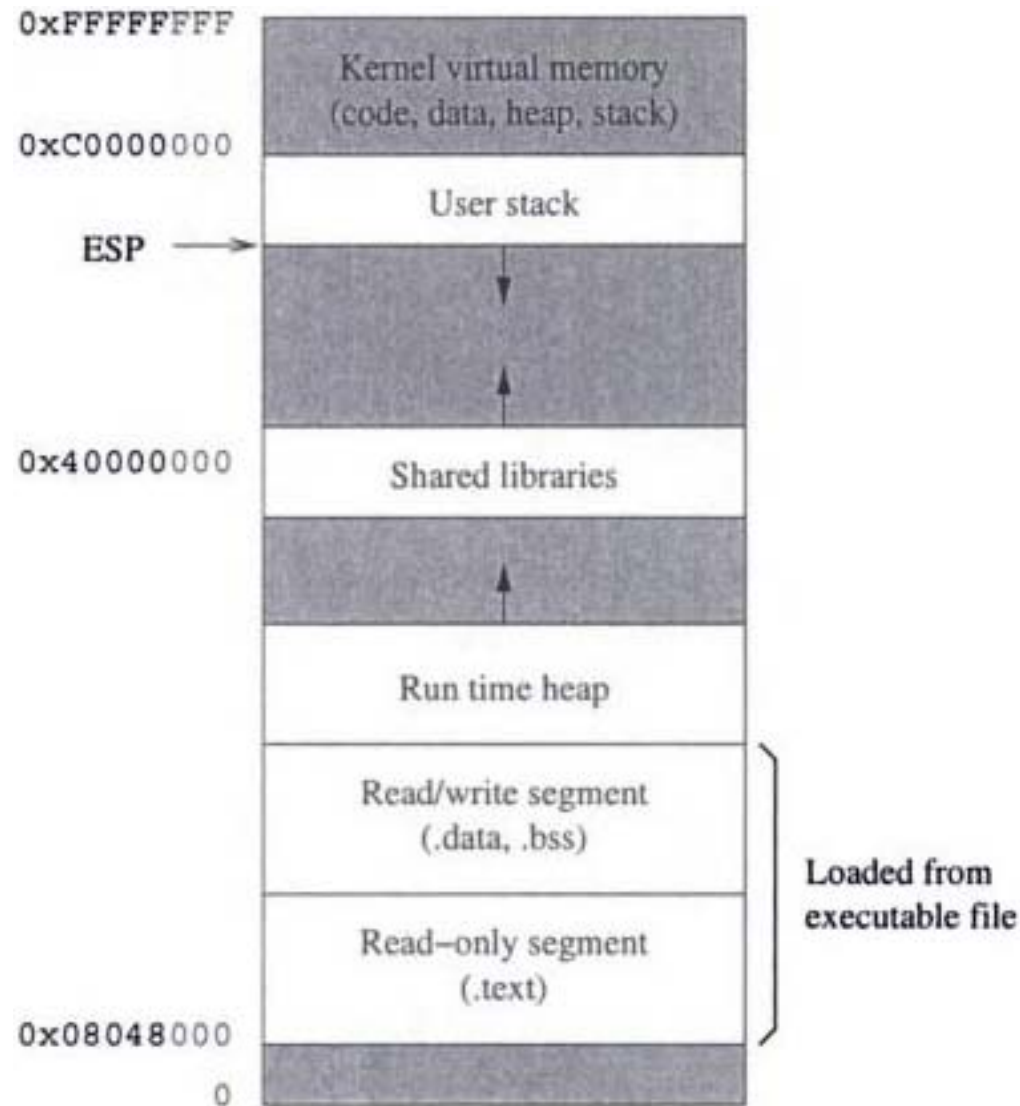
```
nasm -f elf hello.asm
```

```
ld -s -o hello hello.o
```

To run the program, type:

```
./hello
```

# Linux Program Memory



# Parts of the Program

**.data** section contains initialized variables.

Constants can be defined using **equ**

**db, dw, dd, dq, dt** directives can be used.

```
section .data
message: db 'Hello world!'
msglength: equ 12
bufferize: dw 1024
```

# Parts of the Program

**.bss** section contains uninitialized/zeroed variables. **resb**, **resw**, **resd**, **resq**, **rest** directives can be used.

```
section .bss
filename:      resb 255    ;reserve 255 bytes
number:       resb 1      ;reserve 1 byte
num2:         resw 1      ;reserve 1 word
realnum:      resq 10     ;reserve an array
                           of 10 reals
```

# Parts of the Program

**.text** section contains the program in machine language.

declaration `global _start` tells the kernel where program execution begins.

```
section .text  
    global _start
```

```
start:  
    pop ebx  
    .  
    .  
    .
```

# Linux System Calls

- Linux system calls are called in exactly the same way as DOS system calls.
- You put the system call number in EAX (we're dealing with 32-bit registers here, remember)
- You set up the arguments to the system call in EBX, ECX, etc.
- You call the relevant interrupt (for DOS, 21h; for Linux, 80h)
- The result is usually returned in EAX
- There are six registers that are used for the arguments that the system call takes. The first argument goes in EBX, the second in ECX, then EDX, ESI, EDI, and finally EBP, if there are so many. If there are more than six arguments, EBX must contain the memory location where the list of arguments is stored - but don't worry about this because it's unlikely that you'll use a syscall with more than six arguments. The wonderful thing about this scheme is that Linux uses it consistently – all system calls are designed this way, there are no confusing exceptions.



# System Call for Program Exit

```
mov    eax, 1    ;the exit syscall number
mov    ebx, 0    ;exit code of 0
int    80h
```

List of Linux Syscalls (There are totally 190 syscalls):

[http://docs.cs.up.ac.za/programming/asm/derick\\_tut/syscalls.html](http://docs.cs.up.ac.za/programming/asm/derick_tut/syscalls.html)

# Some Linux Syscalls

%eax	Name	Source	%ebx	%ecx	%edx
1	sys_exit	<a href="#">kernel/exit.c</a>	int	-	-
2	sys_fork	<a href="#">arch/i386/kernel/process.c</a>	<a href="#">struct pt_regs</a>	-	-
3	sys_read	<a href="#">fs/read_write.c</a>	unsigned int	char *	<a href="#">size_t</a>
4	sys_write	<a href="#">fs/read_write.c</a>	unsigned int	const char *	<a href="#">size_t</a>
5	sys_open	<a href="#">fs/open.c</a>	const char *	int	int
6	sys_close	<a href="#">fs/open.c</a>	unsigned int	-	-
7	sys_waitpid	<a href="#">kernel/exit.c</a>	pid_t	unsigned int *	int
8	sys_creat	<a href="#">fs/open.c</a>	const char *	int	-
9	sys_link	<a href="#">fs/namei.c</a>	const char *	const char *	-
10	sys_unlink	<a href="#">fs/namei.c</a>	const char *	-	-
11	sys_execve	<a href="#">arch/i386/kernel/process.c</a>	<a href="#">struct pt_regs</a>	-	-
12	sys_chdir	<a href="#">fs/open.c</a>	const char *	-	-
13	sys_time	<a href="#">kernel/time.c</a>	int *	-	-

# Hello World!

```
section .data
hello: db 'Hello world!',0xA ; 0xA linefeed character
helloLen: equ $-hello

section .text
    global _start

_start:
    mov eax,4 ; The system call for write (sys_write)
    mov ebx,1 ; File descriptor 1 - standard output
    mov ecx,hello ; Put the offset of hello in ecx
    mov edx,helloLen

    mov eax,1 ;
    mov ebx,0 ; Exit with return code of 0 (no error)
    int 80h
```

# Initial Stack Layout

<b>argc</b>	[dword] argument counter (integer)
<b>argv[0]</b>	[dword] program name (pointer)
<b>argv[1]</b>	[dword] program args (pointers)
<b>...</b>	
<b>argv[argc-1]</b>	
<b>NULL</b>	[dword] end of args (integer)

**argc and argv[0] are always present**

# Command Line Arguments

- Getting the arguments in C:

```
int main(int argc, char **argv)
{
    while(argc--)
        printf("%s\n", argv[argc]);

    return 0;
}
```

# Command Line Arguments

Popping:

```
pop eax ;get argument counter  
pop ebx ;get our name (argv[0])
```

.arg:

```
pop ecx ;pop all arguments  
test ecx,ecx  
jnz .arg
```

# Command Line Arguments

Direct accessing:

```
pop eax ;get argument counter
```

```
pop esi ;start of arguments
```

```
mov edi,[esp+eax*4] ;end of args
```

# “Procedures” and Jumping

This is a procedure:

```
fileWrite:      mov  eax,4
                 mov  ebx,[filedesc]
                 mov  ecx,stuffToWrite
                 mov  edx,[stuffLen]
                 int  80h
                 ret
```

When coding, procedures and labeling have no difference. The only thing is that:

- You must call a procedure! Otherwise you must jump to the label.



# Example: Find the sum of 4 integers

```
SECTION .data
```

```
x:
```

```
    dd 1
```

```
    dd 5
```

```
    dd 2
```

```
    dd 18
```

```
sum:
```

```
    dd 0
```

```
SECTION .text
```

```
    mov eax,4 ; EAX will serve as a counter
```

```
    mov ebx,0 ; EBX will store the sum
```

```
    mov ecx, x ; ECX will point to the current
```

```
top:
```

```
    add ebx, [ecx]
```

```
    add ecx,4 ; move pointer to next element
```

```
    dec eax ; decrement counter
```

```
    jnz top ; if counter not 0, then loop again
```

```
done:
```

```
    mov [sum],ebx ; done, store result in "sum"
```

# Hello World - 2

- Write a program namely **writef**
- The program takes some arguments from the command line.
- The first arg is the filename, the rest of them are strings to be written into the file.

Usage:

```
./writef output.txt Hello World!
```

# Hello World - 2

section .data

```
        hello    db          'Hello, world!',10    ; Our dear string
        helloLen equ        $ - hello              ; Length of our dear string
```

section .bss

```
address: resw    1
```

section .text

```
global _start
```

\_start:

```
        pop      ebx                ; argc (argument count)
        pop      ebx                ; argv[0] (argument 0, the program name)
        pop      ebx                ; The first real arg, a filename
```

```
        mov      eax,8              ; The syscall number for creat() (we already
have the filename in ebx)
        mov      ecx,00644Q        ; Read/write permissions in octal (rw_rw_rw_)
        int      80h               ; Call the kernel
                                           ; Now we have a file descriptor in eax
```

```
        test     eax,eax            ; Lets make sure the file descriptor is valid
        js       skipWrite ;
```

# Hello World – 2 (cont'd)

filewrite "procedure"

call       fileWrite

skipWrite:

mov       ebx,eax               ; If there was an error, save the errno in ebx  
mov       eax,1                ; Put the exit syscall number in eax  
int       80h                 ; Bail out

fileWrite:

mov       ebx,eax               ; sys\_creat returned file descriptor into eax, now  
move into ebx  
mov       eax,4                ; sys\_write  
                              ; ebx is already set up  
mov       ecx,hello ; We are putting the ADDRESS of hello in ecx  
mov       edx,helloLen         ; This is the VALUE of helloLen because it's a  
constant (defined with equ)

int       80h

mov       eax,6                ; sys\_close (ebx already contains file descriptor)  
int       80h  
ret

# main() { }

- Consider the following program “1.c”:

```
main()  
{  
}
```

- gcc -S -masm=intel 1.c

- Output “1.s” will be:

```
section .text  
global main
```

```
main:
```

```
    push    ebp  
    mov     ebp, esp  
    pop     ebp  
    ret
```

# Passing parameters to functions

- Consider the following program “2.c”:

```
func(int a, int b, int c)
{
    return a+b+c;
}
main()
{
    int x, y=3;
    x=func(y,2,1);
}
```

- gcc -S -O0 -masm=intel 2.c
- Output “2.s” will be:


# Passing parameters to functions

func:

```
push ebp
mov  ebp, esp
mov  eax, DWORD PTR [ebp+8]
add  eax, DWORD PTR [ebp+12]
add  eax, DWORD PTR [ebp+16]
pop  ebp
ret
```

main:

```
push ebp
mov  ebp, esp
sub  esp, 28
mov  DWORD PTR [ebp-8], 3
mov  DWORD PTR [esp+8], 1
mov  DWORD PTR [esp+4], 2
mov  eax, DWORD PTR [ebp-8]
mov  DWORD PTR [esp], eax
call func
mov  DWORD PTR [ebp-4], eax
leave
ret
```



# enter and leave instructions

enter and leave instructions are used to reserve a stack space and releases it. They are equivalent to:

enter n, 0=

push ebp

mov ebp, esp

sub esp, n

leave=

mov esp, ebp

pop ebp



# Global and local variables

```
int var=5;
main()
{
    int a=7;
    int b=9;
    int c=11;
    int d=13;
}
```

***gcc -S -O0 -masm=intel glblcl.c***

```
section .data
    var: dw 5
section .text
global main
```

```
main:
    enter 16,0
    mov word [ebp-4], 7
    mov word [ebp-8], 9
    mov word [ebp-12], 11
    mov word [ebp-16], 13

    leave
    mov eax,1
    mov ebx,0
    int 80h
```

# Compiling/Linking C and Assembly programs - 1

```
extern void puts2(char *, int);
```

```
main()  
{  
    char s[50];  
    sprintf(s,"Hello World 3\n");  
  
    puts2(s,strlen(s));  
}
```

***gcc -c ex4.c***

# Compiling/Linking

## C and Assembly programs - 2

```
section .text  
global puts2
```

```
puts2:  
    push ebp  
    mov ebp, esp  
    mov eax, 4  
    mov ebx, 1  
    mov ecx, dword [ebp+8]  
    mov edx, dword [ebp+12]  
  
    int 80h  
  
    pop ebp  
    ret
```

***nasm -f elf ex5.asm***

# Compiling/Linking C and Assembly programs - 3

```
ayav@ub910TA:~$ gcc -c ex4.c
```

```
ayav@ub910TA:~$ nasm -f elf ex5.asm
```

```
ayav@ub910TA:~$ gcc -o example ex4.o ex5.o
```

```
ayav@ub910TA:~$ ./example
```

**Hello World 3**

```
ayav@ub910TA:~$
```

# Why do we still need to write Assembly programs?

```
#include<stdio.h>
int sub(int a, int b)
{
    return a-b;
}
main()
{
    int x=11;
    x=sub(x,1);
    printf("%u\n",x);
}
```

**gcc -S exp6.c**



sub:

```
.
.
.
push ebp
mov  ebp,esp
mov  eax,DWORD PTR [ebp+12]
mov  edx,DWORD PTR [ebp+8]
mov  ecx,edx
mov  ecx,eax
mov  eax,ecx
pop  ebp
ret
.
.
.
```

# Why do we still need to write Assembly programs?

instead, we can write  
exp7.asm to implement function *sub*:

```
.  
.sub:  
push ebp  
mov ebp,esp  
mov eax,DWORD PTR [ebp+12]  
mov edx,DWORD PTR [ebp+8]  
mov ecx,edx  
sub ecx,eax  
mov eax,ecx  
pop ebp  
ret  
.  
.  
.
```

```
section .text  
global sub  
sub:
```

```
push ebp  
mov ebp,esp  
mov eax,dword [ebp+8]  
sub eax,dword [ebp+8]  
pop ebp  
ret
```

Return  
value



**Our code is shorter than the one gcc produced.**

Note: how return value is passed to the caller program

# Pointer arguments

```
#include<stdio.h>
```

```
void signal(int *s)
{
    (*s)++;
}
```

```
main()
{
    int sem=0;
    signal(&sem);
    printf("%u\n",sem);
}
```

```
section .text
```

```
global signal
```

```
signal:
```

```
    push ebp
```

```
    mov ebp,esp
```

```
    mov eax,dword [ebp+8]
```

```
    mov edx, [eax]
```

```
    inc edx
```

```
    mov [eax],edx
```

```
    pop ebp
```

```
    ret
```



# Inline Assembly



# AT&T vs. Intel Syntax

- GNU C compiler uses AT&T syntax

Intel Code	AT&T Code
mov eax, 1	movl \$1, %eax
mov ebx, 0ffh	movl \$0xff, %ebx
int 80h	int \$0x80
mov ebx, eax	movl %eax, %ebx
mov eax, [ecx]	movl (%ecx), %eax
mov eax, [ebx&plus;3]	movl 3(%ebx), %eax
mov eax, [ebx+20h]	movl 0x20(%ebx), %eax
add eax, [ebx+ecx*2h]	addl (%ebx, %ecx, 0x2), %eax
lea eax, [ebx+ecx]	leal (%ebx, %ecx), %eax
sub eax, [ebx+ecx*4h-20h]	subl -0x20(%ebx, %ecx, 0x4), %eax

# Basic inline assembly

```
asm( "movl  %ecx %eax" );
```

```
/* moves the contents of ecx to eax */
```

```
__asm__( "movb  %bh  ( %eax ) " );
```

```
/*moves the byte from bh to the memory pointed  
by eax */
```

# Basic inline assembly

```
__asm__ ( "movl %eax, %ebx\n\t"  
          "movl $56, %esi\n\t"  
          "movl %ecx,  
$label(%edx,%ebx,$4)\n\t"  
          "movb %ah, (%ebx)" );
```

We touch some registers and exit without fixing them.  
This may cause big trouble

# Extended inline assembly

```
asm ( assembler template  
    : output operands /* optional */  
    : input operands /* optional */  
    : list of clobbered registers /* optional */  
    );
```

# Extended inline assembly

```
int a=10, b;  
asm ( "    movl %1, %%eax;  
        movl %%eax, %0;"  
      : "=r"(b) /* output */  
      : "r"(a) /* input */  
      : "%eax" /* clobbered register */  
    );
```

# Some Recipes - I

```
int main(void) {  
  
    int foo = 10, bar = 15;  
    __asm__ __volatile__(  
        "addl %%ebx,%%eax"  
        : "=a" (foo) : "a" (foo), "b" (bar) );  
  
    printf("foo+bar=%d\n", foo);  
  
    return 0; }
```

# Some Recipes – II

## Move Block Macro

```
#define mov_blk(src, dest, numwords) \
__asm__ __volatile__ ( \
    "cld\n\t" \
    "rep\n\t" \
    "movsl" \
    : : "S" (src), "D" (dest), "c" (numwords) \
    : "%ecx", "%esi", "%edi" \
)
```

# Some Recipes – III

## Exit System Call

```
asm( "movl $1,%%eax; /* SYS_exit is 1 */  
xorl %%ebx,%%ebx;  
    /* Argument is in ebx, it is 0 */  
int $0x80" /* Enter kernel mode */  
);
```



# Some Recipes – IV

## Exit System Call

```
int reason;
```

```
asm( "movl $1,%%eax; /* SYS_exit is 1 */  
int $0x80" /* Enter kernel mode  
: : "b" (reason)  
: "%eax" "%ebx"  
);
```